


a V2 multimodal inference architecture

Image-dominant gated fusion keeps ultrasound evidence primary while incorporating exam text.



Training scale: 10,587 patients / 196,255 images | 5-fold CV | Fold 3 EMA pat_acc=0.8600